

Project 3: The Kernel Crypto API

Group 10-03: Behnam Saeedi, zhaoheng wang, Levi Willmeth
CS444: Operating systems II



Spring 2017

Abstract

CONTENTS

1	Design	3
1.1	Description of concept	3
1.2	Application	3
1.3	Algorithm	3
1.4	Pseudo Codes	3
2	Q&A	3
2.1	Questions	3
2.1.1	What do you think the main point of this assignment is?	3
2.1.2	How did you personally approach the problem? Design decisions, algorithm, etc.	4
2.1.3	How did you ensure your solution was correct? Testing details, for instance.	4
2.1.4	What did you learn?	4
3	Logs	4
3.1	Work-flow and scheduling	4
3.2	Work log	4
3.3	Git commit log	5
4	Code	6
4.1	rd2.c	6
4.2	Kconfig	9
4.3	Makefile	19
4.4	Patch	20

1 DESIGN

1.1 Description of concept

Write a disk device driver module for linux-yocto-3.14 that can perform encryption based on a key passed as a parameter. Test our solution. Submit it as a Linux patch. Don't make any mistakes (test our submission).

1.2 Application

We began by reading the textbook Linux Device Drivers, Third Edition. Chapter 16 guided us towards modifying the sbull.c file provided by the publisher. We learned how to build and run the module while reading chapter 2. We follow the rules on the chapter 2 and modify the sbull.c file. After that, we set the driver to build as module by typing make menuconfig and build the kernel to generate the rd2.ko. We scp the rd2.ko into QEMU and (test)

1.3 Algorithm

Allocate a block of memory. Any time we write into that memory, apply the crypto algorithm and key to the data before writing. Any time we read from that memory, apply the crypto algorithm and key before returning the data.

1.4 Pseudo Codes

```
Allocate block of memory per request
make crypto cipher using key
if writing to memory:
    memset = (start,0,end) //set the memory to the data
    data = encrypt(data)
else:
    data = decrypt(data)
return data
```

2 Q&A

2.1 Questions

2.1.1 What do you think the main point of this assignment is?

The main point of this assignment was to force us to get up close and personal with several new areas of Linux kernel programming: modifying device drivers to fit our needs, introducing applying practical encryption on the block level, and packaging our solution as a Linux module. Each of these tasks had it's own challenges and obstacles to overcome. Each task relied on the previous, meaning that we needed to understand both the steps we had achieved so far, and what we were trying to accomplish with our next step.

2.1.2 How did you personally approach the problem? Design decisions, algorithm, etc.

We began by trying to understand the problem. What was it really asking us to do? How could we accomplish those steps? We leaned heavily on the class textbook, *Linux Device Drivers, Third Edition*. Chapter 16 led us through writing a device driver based on `sbull.c` and chapter 2 helped us to understand how to package our solution as a module. Most of the steps in between were solved with a combination of group discussion, moments of inspiration, and plenty of searching the Internet for similar errors. For example, we got stuck for several hours on a compilation error in the textbook's example file `sbull.c`, before finally realizing it relied on a different version of the kernel's `bio.h`.

2.1.3 How did you ensure your solution was correct? Testing details, for instance.

We knew that our module should be performing encryption. So how could we test something that should be performing encryption. Make a file and write some data to it. If you can read it back, encryption/decryption might be working. But how can we tell for sure? Simple, just disable our module. If the file we made earlier is now gibberish, we know that it was encrypted, and that we were also able to decrypt it while the module was running.

2.1.4 What did you learn?

We learned to give ourselves plenty of time with these assignments. This one had some time-consuming surprises for us. Luckily, we began working on the assignment right away, so we had enough time to work through each problem. The interesting and challenging thing about these assignments is that they have several components, each of which is new for us. In this example, we faced at least 6 new problems: writing a device driver, adding encryption, testing, bundling our solution as a module, then bringing our module into the VM. Many of these problems were not-insignificant challenges that required research, time, and some trial-and-error to both solve and attempt to understand why our solution worked.

3 Logs

3.1 Work-flow and scheduling

We used a different work-flow for different portions of the assignment. We usually worked in a group of 3, but often one or two of us were researching a problem while the others attempted trial-and-error, or probed the problem in an attempt to better understand the details. When we found a solution and one of us understood it, we tried to explain it both to ourselves and the others. Followed by promptly making a git commit or editing a Makefile to recreate the circumstances we used to solve the problem.

For file organization, we create a folder in our team directory for each assignment. If we are attempting to solve a particular problem in more than one way, we create git branches until someone solves it. Then we merge back to master and proceed to the next problem together.

3.2 Work log

Our team met Monday, Wednesday, and Friday from 2-4 pm, and as the deadline approached we also met over the weekend.

3.3 Git commit log

```
commit 29f3058a70efc8de21af6ab6b90c887afbc0e480
Author: Levi Willmeth <levi.willmeth@gmail.com>
Date: 2017-05-22
```

Update Kconfig and Makefile to working versions

```
commit b3b0c6db967cc7ae85ed892111fd05782429b735
Author: Levi Willmeth <levi.willmeth@gmail.com>
Date: 2017-05-22
```

Add rd2 files, update Makefile and Kconfig

```
commit acf53576d3fc0a25128e9f76de41dbda9452b918
Author: Levi Willmeth <levi.willmeth@gmail.com>
Date: 2017-05-22
```

Merge in testing branch

```
commit 0ae768ede536e97d9134e81f439a103411e6648c
Author: wangdaye123 <wangzhao@oregonstate.edu>
Date: 2017-05-15
```

delete unnecessary file

```
commit c5b6ba77d07e996be3b7cbcfcb1272aa8cac283b
Merge: 72a4443 5629be4
Author: wangdaye123 <wangzhao@oregonstate.edu>
Date: 2017-05-15
```

Merge branch '**master**' of <https://github.com/BeNsAeI/CS444>

```
commit 5629be4e2a8ed0d839c7dbe204968e67a2947d1b
Author: wangdaye123 <wangzhao@oregonstate.edu>
Date: 2017-05-15
```

delete the file

```
commit 51d07eec0d70d2b644e2bb8806aa9e70e09dac70
Author: wangdaye123 <wangzhao@oregonstate.edu>
Date: 2017-05-15
```

upload the file need to change

```
commit 490998a4ae14dadded40e12bdf0dc5fccff8060b
Author: wangdaye123 <wangzhao@oregonstate.edu>
Date: 2017-05-15
```

upload requirement of hw3

```
commit 72a4443cd53ea10763b755eb35f566c40f2b2b34
Author: wangdaye123 <wangzhao@oregonstate.edu>
Date: 2017-05-15
```

```

add hw3 requirement and file need to change

commit ef3752d958a125b8f5db236a8ed812c6def6ce08
Author: BeNsAeI <behnam.saeidi@gmail.com>
Date: 2017-05-10

added sstf.cfile

commit f660de28c8c30bb3bf558dffae2ceae88c19fbe6
Author: Behnam Saeedi <saeedib@os-class.engr.oregonstate.edu>
Date: 2017-05-10

backup files

commit 53780e85fd233b72b073bdea2caf50cbcd28d5ea
Author: BeNsAeI <behnam.saeidi@gmail.com>
Date: 2017-05-09

added lecture notes for 9th lecture

```

4 CODE

4.1 rd2.c

```

#include <linux/module.h>
#include <linux/moduleparam.h>
#include <linux/init.h>

#include <linux/kernel.h>
#include <linux/fs.h>
#include <linux/errno.h>
#include <linux/types.h>
#include <linux/vmalloc.h>
#include <linux/genhd.h>
#include <linux/blkdev.h>
#include <linux/hdreg.h>

#include <linux/crypto.h>

MODULE_LICENSE("Dual BSD/GPL");
static char *Version = "1.4";

static struct crypto_cipher *cipher;

static int major_num = 0;
module_param(major_num, int, 0);
static int logical_block_size = 512;
module_param(logical_block_size, int, 0);
static int nsectors = 1024;
module_param(nsectors, int, 0);

#define KEY_SZ 8

```

```

static char* key = "privateKey";
module_param(key, charp, 0);
#define KERNEL_SECTOR_SIZE 512

static struct request_queue *Queue;

static struct rd2_device {
    unsigned long size;
    spinlock_t lock;
    u8 *data;
    struct gendisk *gd;
} Device;

static void rd2_transfer(struct rd2_device *dev, sector_t sector,
    unsigned long nsect, char *buffer, int write) {
    unsigned long offset = sector * logical_block_size;
    unsigned long nbytes = nsect * logical_block_size;
    unsigned int i;

    //printk("[KERNEL DEBUG] rd2: transfer_key > %s\n",key);

    if ((offset + nbytes) > dev->size) {
        printk (KERN_NOTICE "[KERNEL DEBUG] rd2(%s): Beyond-end write (%ld %ld)\n", Version, offset, nb
        return;
    }

    crypto_cipher_clear_flags(cipher, ~0);
    crypto_cipher_setkey(cipher, key, strlen(key));

    if(write) {
        for(i = 0; i < nbytes; i += crypto_cipher_blocksize(cipher)) {
            memset(dev->data + offset + i, 0, crypto_cipher_blocksize(cipher));
            crypto_cipher_encrypt_one(cipher, dev->data + offset + i, buffer + i);
        }
    } else {
        for(i = 0; i < nbytes; i += crypto_cipher_blocksize(cipher)) {
            crypto_cipher_decrypt_one(cipher, buffer + i, dev->data + offset + i);
        }
    }
}

static void rd2_request(struct request_queue *q) {
    struct request *req;

    req = blk_fetch_request(q);
    while (req != NULL) {
        if (req == NULL || (req->cmd_type != REQ_TYPE_FS)) {
            __blk_end_request_all(req, -EIO);
            continue;
        }
        rd2_transfer(&Device, blk_rq_pos(req), blk_rq_cur_sectors(req),
            bio_data(req->bio), rq_data_dir(req));

        if ( ! __blk_end_request_cur(req, 0) ) {
            req = blk_fetch_request(q);

```

```

}
}
}

int rd2_getgeo(struct block_device * block_device, struct hd_geometry * geo) {
long size;

size = Device.size * (logical_block_size / KERNEL_SECTOR_SIZE);
geo->cylinders = (size & ~0x3f) >> 6;
geo->heads = 4;
geo->sectors = 16;
geo->start = 0;
return 0;
}

static struct block_device_operations rd2_ops = {
.owner = THIS_MODULE,
.getgeo = rd2_getgeo
};

static int __init rd2_init(void) {
Device.size = nsectors * logical_block_size;
spin_lock_init(&Device.lock);
Device.data = vmalloc(Device.size);
if (Device.data == NULL)
return -ENOMEM;
/*
 * Get a request queue.
 */
Queue = blk_init_queue(rd2_request, &Device.lock);
if (Queue == NULL)
goto out;
blk_queue_logical_block_size(Queue, logical_block_size);
/*
 * Get registered.
 */
major_num = register_blkdev(major_num, "rd2");
if (major_num < 0) {
goto out;
}

/* allocating cipher */
cipher = crypto_alloc_cipher("aes",0,0);
if(!cipher){
goto out;
}

printk("[KERNEL DEBUG] rd2: enc_key > %s\n",key);
crypto_cipher_setkey(cipher,key,strlen(key));

/*
 * And the gendisk structure.
 */
Device.gd = alloc_disk(16);
if (!Device.gd)

```



```

goto out_unregister;
Device.gd->major = major_num;
Device.gd->first_minor = 0;
Device.gd->fops = &rd2_ops;
Device.gd->private_data = &Device;
strcpy(Device.gd->disk_name, "rd20");
set_capacity(Device.gd, nsectors);
Device.gd->queue = Queue;
add_disk(Device.gd);
return 0;

out_unregister:
unregister_blkdev(major_num, "rd2");
out:
vfree(Device.data);
return -ENOMEM;
}

static void __exit rd2_exit(void){

crypto_free_cipher(cipher);

del_gendisk(Device.gd);
put_disk(Device.gd);
unregister_blkdev(major_num, "rd2");
blk_cleanup_queue(Queue);
}

module_init(rd2_init);
module_exit(rd2_exit);

```

4.2 Kconfig

```

#
# Block device driver configuration
#

menuconfig BLK_DEV
    bool "Block devices"
    depends on BLOCK
    default y
    ---help---
    Say Y here to get to see options for various different block device
    drivers. This option alone does not add any kernel code.

    If you say N, all options in this submenu will be skipped and disabled;
    only do this if you know what you are doing.

if BLK_DEV

config BLK_DEV_RD2
    tristate "RD2 DRIVER TEST 10-03"

config BLK_DEV_NULL_BLK

```

```

        tristate "Null test block driver"

config BLK_DEV_FD
    tristate "Normal floppy disk support"
    depends on ARCH_MAY_HAVE_PC_FDC
    ---help---
        If you want to use the floppy disk drive(s) of your PC under Linux,
        say Y. Information about this driver, especially important for IBM
        Thinkpad users, is contained in
        <file:Documentation/blockdev/floppy.txt>.
        That file also contains the location of the Floppy driver FAQ as
        well as location of the fdutils package used to configure additional
        parameters of the driver at run time.

        To compile this driver as a module, choose M here: the
        module will be called floppy.

config AMIGA_FLOPPY
    tristate "Amiga floppy support"
    depends on AMIGA

config ATARI_FLOPPY
    tristate "Atari floppy support"
    depends on ATARI

config MAC_FLOPPY
    tristate "Support for PowerMac floppy"
    depends on PPC_PMAC && !PPC_PMAC64
    help
        If you have a SWIM-3 (Super Woz Integrated Machine 3; from Apple)
        floppy controller, say Y here. Most commonly found in PowerMacs.

config BLK_DEV_SWIM
    tristate "Support for SWIM Macintosh floppy"
    depends on M68K && MAC
    help
        You should select this option if you want floppy support
        and you don't have a II, IIx, Q900, Q950 or AV series.

config AMIGA_Z2RAM
    tristate "Amiga Zorro II ramdisk support"
    depends on ZORRO
    help
        This enables support for using Chip RAM and Zorro II RAM as a
        ramdisk or as a swap partition. Say Y if you want to include this
        driver in the kernel.

        To compile this driver as a module, choose M here: the
        module will be called z2ram.

config GDRAM
    tristate "SEGA Dreamcast GD-ROM drive"
    depends on SH_DREAMCAST
    help
        A standard SEGA Dreamcast comes with a modified CD ROM drive called a

```

"GD-ROM" by SEGA to signify it is capable of reading special disks with up to 1 GB of data. This drive will also read standard CD ROM disks. Select this option to access any disks in your GD ROM drive. Most users will want to say "Y" here.
You can also build this as a module which will be called gdrom.

config PARIDE

tristate "Parallel port IDE device support"
depends on PARPORT_PC

---help---

There are many external CD-ROM and disk devices that connect through your computer's parallel port. Most of them are actually IDE devices using a parallel port IDE adapter. This option enables the PARIDE subsystem which contains drivers for many of these external drives. Read <file:Documentation/blockdev/paride.txt> for more information.

If you have said Y to the "Parallel-port support" configuration option, you may share a single port between your printer and other parallel port devices. Answer Y to build PARIDE support into your kernel, or M if you would like to build it as a loadable module. If your parallel port support is in a loadable module, you must build PARIDE as a module. If you built PARIDE support into your kernel, you may still build the individual protocol modules and high-level drivers as loadable modules. If you build this support as a module, it will be called paride.

To use the PARIDE support, you must say Y or M here and also to at least one high-level driver (e.g. "Parallel port IDE disks", "Parallel port ATAPI CD-ROMs", "Parallel port ATAPI disks" etc.) and to at least one protocol driver (e.g. "ATEN EH-100 protocol", "MicroSolutions backpack protocol", "DataStor Commuter protocol" etc.).

source "drivers/block/paride/Kconfig"

source "drivers/block/mtip32xx/Kconfig"

source "drivers/block/zram/Kconfig"

config BLK_CPQ_DA

tristate "Compaq SMART2 support"
depends on PCI && VIRT_TO_BUS && 0

help

This is the driver for Compaq Smart Array controllers. Everyone using these boards should say Y here. See the file <file:Documentation/blockdev/cpqarray.txt> for the current list of boards supported by this driver, and for further information on the use of this driver.

config BLK_CPQ_CISS_DA

tristate "Compaq Smart Array 5xxx support"
depends on PCI

select CHECK_SIGNATURE

help

This is the driver for Compaq Smart Array 5xxx controllers.

Everyone using these boards should say Y here.
 See <file:Documentation/blockdev/cciss.txt> **for** the current list of
 boards supported by this driver, and **for** further information
 on the use of this driver.

```
config CISS SCSI_TAPE
bool "SCSI tape drive support for Smart Array 5xxx"
depends on BLK_CPQ_CISS_DA && PROC_FS
depends on SCSI=y || SCSI=BLK_CPQ_CISS_DA
help
  When enabled (Y), this option allows SCSI tape drives and SCSI medium
  changers (tape robots) to be accessed via a Compaq 5xxx array
  controller. (See <file:Documentation/blockdev/cciss.txt> for more details.)

  "SCSI support" and "SCSI tape support" must also be enabled for this
  option to work.

  When this option is disabled (N), the SCSI portion of the driver
  is not compiled.
```

```
config BLK_DEV_DAC960
tristate "Mylex DAC960/DAC1100 PCI RAID Controller support"
depends on PCI
help
  This driver adds support for the Mylex DAC960, AcceleRAID, and
  eXtremeRAID PCI RAID controllers. See the file
  <file:Documentation/blockdev/README.DAC960> for further information
  about this driver.

  To compile this driver as a module, choose M here: the
  module will be called DAC960.
```

```
config BLK_DEV_UMEM
tristate "Micro Memory MM5415 Battery Backed RAM support"
depends on PCI
---help---
  Saying Y here will include support for the MM5415 family of
  battery backed (Non-volatile) RAM cards.
  <http://www.umem.com/>

  The cards appear as block devices that can be partitioned into
  as many as 15 partitions.

  To compile this driver as a module, choose M here: the
  module will be called umem.

  The umem driver has not yet been allocated a MAJOR number, so
  one is chosen dynamically.
```

```
config BLK_DEV_UBD
bool "Virtual block device"
depends on UML
---help---
  The User-Mode Linux port includes a driver called UBD which will let
  you access arbitrary files on the host computer as block devices.
```

Unless you know that you **do** not need such virtual block devices say Y here.

```
config BLK_DEV_UBD_SYNC
    bool "Always do synchronous disk IO for UBD"
    depends on BLK_DEV_UBD
    ---help---
        Writes to the virtual block device are not immediately written to the
        host's disk; this may cause problems if, for example, the User-Mode
        Linux 'Virtual Machine' uses a journalling filesystem and the host
        computer crashes.

        Synchronous operation (i.e. always writing data to the host's disk
        immediately) is configurable on a per-UBD basis by using a special
        kernel command line option. Alternatively, you can say Y here to
        turn on synchronous operation by default for all block devices.

        If you're running a journalling file system (like reiserfs, for
        example) in your virtual machine, you will want to say Y here. If
        you care for the safety of the data in your virtual machine, Y is a
        wise choice too. In all other cases (for example, if you're just
        playing around with User-Mode Linux) you can choose N.
```

```
config BLK_DEV_COW_COMMON
    bool
    default BLK_DEV_UBD
```

```
config BLK_DEV_LOOP
    tristate "Loopback device support"
    ---help---
        Saying Y here will allow you to use a regular file as a block
        device; you can then create a file system on that block device and
        mount it just as you would mount other block devices such as hard
        drive partitions, CD-ROM drives or floppy drives. The loop devices
        are block special device files with major number 7 and typically
        called /dev/loop0, /dev/loop1 etc.

        This is useful if you want to check an ISO 9660 file system before
        burning the CD, or if you want to use floppy images without first
        writing them to floppy. Furthermore, some Linux distributions avoid
        the need for a dedicated Linux partition by keeping their complete
        root file system inside a DOS FAT file using this loop device
        driver.
```

To use the loop device, you need the losetup utility, found in the util-linux package, see [<ftp://ftp.kernel.org/pub/linux/utils/util-linux/>](http://ftp.kernel.org/pub/linux/utils/util-linux/).

The loop device driver can also be used to **"hide"** a file system in a disk partition, floppy, or regular file, either using encryption (scrambling the data) or steganography (hiding the data in the low bits of, say, a sound file). This is also safe **if** the file resides on a remote file server.

There are several ways of encrypting disks. Some of these require

kernel patches. The vanilla kernel offers the cryptoloop option and a Device Mapper target (which is superior, as it supports all file systems). If you want to use the cryptoloop, say Y to both LOOP and CRYPTOLOOP, and make sure you have a recent (version 2.12 or later) version of util-linux. Additionally, be aware that the cryptoloop is not safe **for** storing journaled filesystems.

Note that this loop device has nothing to **do** with the loopback device used **for** network connections from the machine to itself.

To compile this driver as a module, choose M here: the module will be called loop.

Most users will answer N here.

```
config BLK_DEV_LOOP_MIN_COUNT
    int "Number of loop devices to pre-create at init time"
    depends on BLK_DEV_LOOP
    default 8
    help
        Static number of loop devices to be unconditionally pre-created
        at init time.

        This default value can be overwritten on the kernel command
        line or with module-parameter loop.max_loop.

        The historic default is 8. If a late 2011 version of losetup(8)
        is used, it can be set to 0, since needed loop devices can be
        dynamically allocated with the /dev/loop-control interface.

config BLK_DEV_CRYPTOLOOP
    tristate "Cryptoloop Support"
    select CRYPTO
    select CRYPTO_CBC
    depends on BLK_DEV_LOOP
    ---help---
        Say Y here if you want to be able to use the ciphers that are
        provided by the CryptoAPI as loop transformation. This might be
        used as hard disk encryption.

        WARNING: This device is not safe for journaled file systems like
        ext3 or Reiserfs. Please use the Device Mapper crypto module
        instead, which can be configured to be on-disk compatible with the
        cryptoloop device.

source "drivers/block/drbd/Kconfig"

config BLK_DEV_NBD
    tristate "Network block device support"
    depends on NET
    ---help---
        Saying Y here will allow your computer to be a client for network
        block devices, i.e. it will be able to use block devices exported by
        servers (mount file systems on them etc.). Communication between
        client and server works over TCP/IP networking, but to the client
```

program this is hidden: it looks like a regular `local` file access to a block device special file such as `/dev/nd0`.

Network block devices also allows you to run a block-device in userland (making server and client physically the same computer, communicating using the loopback network device).

Read `<file:Documentation/blockdev/nbd.txt>` `for` more information, especially about where to find the server code, which runs in user space and does not need special kernel support.

Note that this has nothing to `do` with the network file systems NFS or Coda; you can say `N` here even `if` you intend to use NFS or Coda.

To compile this driver as a module, choose `M` here: the module will be called `nbd`.

If unsure, say `N`.

```
config BLK_DEV_NVME
    tristate "NVM Express block device"
    depends on PCI
    ---help---
    The NVM Express driver is for solid state drives directly
    connected to the PCI or PCI Express bus. If you know you
    don't have one of these, it is safe to answer N.

    To compile this driver as a module, choose M here: the
    module will be called nvme.
```

```
config BLK_DEV_SKD
    tristate "STEC S1120 Block Driver"
    depends on PCI
    depends on 64BIT
    ---help---
    Saying Y or M here will enable support for the
    STEC, Inc. S1120 PCIe SSD.

    Use device /dev/skd$N and /dev/skd$Np$M.
```

```
config BLK_DEV_OSD
    tristate "OSD object-as-blkdev support"
    depends on SCSI_OSD_ULD
    ---help---
    Saying Y or M here will allow the exporting of a single SCSI
    OSD (object-based storage) object as a Linux block device.

    For example, if you create a 2G object on an OSD device,
    you can then use this module to present that 2G object as
    a Linux block device.

    To compile this driver as a module, choose M here: the
    module will be called osdbl.

    If unsure, say N.
```

```

config BLK_DEV_SX8
    tristate "Promise SATA SX8 support"
    depends on PCI
    ---help---
        Saying Y or M here will enable support for the
        Promise SATA SX8 controllers.

        Use devices /dev/sx8/$N and /dev/sx8/$Np$M.

config BLK_DEV_RAM
    tristate "RAM block device support"
    ---help---
        Saying Y here will allow you to use a portion of your RAM memory as
        a block device, so that you can make file systems on it, read and
        write to it and do all the other things that you can do with normal
        block devices (such as hard drives). It is usually used to load and
        store a copy of a minimal root file system off of a floppy into RAM
        during the initial install of Linux.

        Note that the kernel command line option "ramdisk=XX" is now obsolete.
        For details, read <file:Documentation/blockdev/ramdisk.txt>.

        To compile this driver as a module, choose M here: the
        module will be called brd. An alias "rd" has been defined
        for historical reasons.

        Most normal users won't need the RAM disk functionality, and can
        thus say N here.

config BLK_DEV_RAM_COUNT
    int "Default number of RAM disks"
    default "16"
    depends on BLK_DEV_RAM
    help
        The default value is 16 RAM disks. Change this if you know what you
        are doing. If you boot from a filesystem that needs to be extracted
        in memory, you will need at least one RAM disk (e.g. root on cramfs).

config BLK_DEV_RAM_SIZE
    int "Default RAM disk size (kbytes)"
    depends on BLK_DEV_RAM
    default "4096"
    help
        The default value is 4096 kilobytes. Only change this if you know
        what you are doing.

config BLK_DEV_XIP
    bool "Support XIP filesystems on RAM block device"
    depends on BLK_DEV_RAM
    default n
    help
        Support XIP filesystems (such as ext2 with XIP support on) on
        top of block ram device. This will slightly enlarge the kernel, and
        will prevent RAM block device backing store memory from being

```


allocated from highmem (only a problem **for** highmem systems).

```
config CDROM_PKTCDVD
    tristate "Packet writing on CD/DVD media"
    depends on !UML
    help
        If you have a CDROM/DVD drive that supports packet writing, say
        Y to include support. It should work with any MMC/Mt Fuji
        compliant ATAPI or SCSI drive, which is just about any newer
        DVD/CD writer.

        Currently only writing to CD-RW, DVD-RW, DVD+RW and DVD-RAM discs
        is possible.
        DVD-RW disks must be in restricted overwrite mode.

        See the file <file:Documentation/cdrom/packet-writing.txt>
        for further information on the use of this driver.

        To compile this driver as a module, choose M here: the
        module will be called pktcdvd.
```

```
config CDROM_PKTCDVD_BUFFERS
    int "Free buffers for data gathering"
    depends on CDROM_PKTCDVD
    default "8"
    help
        This controls the maximum number of active concurrent packets. More
        concurrent packets can increase write performance, but also require
        more memory. Each concurrent packet will require approximately 64Kb
        of non-swappable kernel memory, memory which will be allocated when
        a disc is opened for writing.
```

```
config CDROM_PKTCDVD_WCACHE
    bool "Enable write caching"
    depends on CDROM_PKTCDVD
    help
        If enabled, write caching will be set for the CD-R/W device. For now
        this option is dangerous unless the CD-RW media is known good, as we
        don't do deferred write error handling yet.
```

```
config ATA_OVER_ETH
    tristate "ATA over Ethernet support"
    depends on NET
    help
        This driver provides Support for ATA over Ethernet block
        devices like the Coraid EtherDrive (R) Storage Blade.
```

```
config MG_DISK
    tristate "mGine mflash, gflash support"
    depends on ARM && GPIOLIB
    help
        mGine mFlash(gFlash) block device driver
```

```
config MG_DISK_RES
    int "Size of reserved area before MBR"
```

```

depends on MG_DISK
default 0
help
    Define size of reserved area that usually used for boot. Unit is KB.
    All of the block device operation will be taken this value as start
    offset
    Examples:
        1024 => 1 MB

config SUNVDC
    tristate "Sun Virtual Disk Client support"
    depends on SUN_LDOMS
    help
        Support for virtual disk devices as a client under Sun
        Logical Domains.

source "drivers/s390/block/Kconfig"

config XILINX_SYSACE
    tristate "Xilinx SystemACE support"
    depends on 4xx || MICROBLAZE
    help
        Include support for the Xilinx SystemACE CompactFlash interface

config XEN_BLKDEV_FRONTEND
    tristate "Xen virtual block device support"
    depends on XEN
    default y
    select XEN_XENBUS_FRONTEND
    help
        This driver implements the front-end of the Xen virtual
        block device driver. It communicates with a back-end driver
        in another domain which drives the actual block device.

config XEN_BLKDEV_BACKEND
    tristate "Xen block-device backend driver"
    depends on XEN_BACKEND
    help
        The block-device backend driver allows the kernel to export its
        block devices to other guests via a high-performance shared-memory
        interface.

        The corresponding Linux frontend driver is enabled by the
        CONFIG_XEN_BLKDEV_FRONTEND configuration option.

        The backend driver attaches itself to a any block device specified
        in the XenBus configuration. There are no limits to what the block
        device as long as it has a major and minor.

        If you are compiling a kernel to run in a Xen block backend driver
        domain (often this is domain 0) you should say Y here. To
        compile this driver as a module, chose M here: the module
        will be called xen-blkback.

```

```

config VIRTIO_BLK
    tristate "Virtio block driver"
    depends on VIRTIO
    ---help---
        This is the virtual block driver for virtio.  It can be used with
        lguest or QEMU based VMs (like KVM or Xen).  Say Y or M.

config BLK_DEV_HD
    bool "Very old hard disk (MFM/RLL/IDE) driver"
    depends on HAVE_IDE
    depends on !ARM || ARCH_RPC || BROKEN
    help
        This is a very old hard disk driver that lacks the enhanced
        functionality of the newer ones.

        It is required for systems with ancient MFM/RLL/ESDI drives.

        If unsure, say N.

config BLK_DEV_RBD
    tristate "Rados block device (RBD)"
    depends on INET && BLOCK
    select CEPH_LIB
    select LIBCRC32C
    select CRYPTO_AES
    select CRYPTO
    default n
    help
        Say Y here if you want include the Rados block device, which stripes
        a block device over objects stored in the Ceph distributed object
        store.

        More information at http://ceph.newdream.net/.

        If unsure, say N.

config BLK_DEV_RSXX
    tristate "IBM Flash Adapter 900GB Full Height PCIe Device Driver"
    depends on PCI
    help
        Device driver for IBM's high speed PCIe SSD
        storage device: Flash Adapter 900GB Full Height.

        To compile this driver as a module, choose M here: the
        module will be called rsxx.

endif # BLK_DEV

```

4.3 Makefile

```

#
# Makefile for the kernel block device drivers.
#
# 12 June 2000, Christoph Hellwig <hch@infradead.org>

```

```
# Rewritten to use lists instead of if-statements.
#
```

```
obj-$(CONFIG_MAC_FLOPPY)      += swim3.o
obj-$(CONFIG_BLK_DEV_SWIM)    += swim_mod.o
obj-$(CONFIG_BLK_DEV_FD)      += floppy.o
obj-$(CONFIG_AMIGA_FLOPPY)    += amiflop.o
obj-$(CONFIG_PS3_DISK)       += ps3disk.o
obj-$(CONFIG_PS3_VRAM)       += ps3vram.o
obj-$(CONFIG_ATARI_FLOPPY)    += ataflop.o
obj-$(CONFIG_AMIGA_Z2RAM)     += z2ram.o
obj-$(CONFIG_BLK_DEV_RAM)     += brd.o
obj-$(CONFIG_BLK_DEV_LOOP)    += loop.o
obj-$(CONFIG_BLK_CPQ_DA)      += cpqarray.o
obj-$(CONFIG_BLK_CPQ_CISS_DA) += cciss.o
obj-$(CONFIG_BLK_DEV_DAC960)  += DAC960.o
obj-$(CONFIG_XILINX_SYSACE)   += xsysace.o
obj-$(CONFIG_CDROM_PKTCDVD)   += pktcdvd.o
obj-$(CONFIG_MG_DISK)         += mg_disk.o
obj-$(CONFIG_SUNVDC)          += sunvdc.o
obj-$(CONFIG_BLK_DEV_NVME)    += nvme.o
obj-$(CONFIG_BLK_DEV_SKD)     += skd.o
obj-$(CONFIG_BLK_DEV_OSD)     += osdblk.o

obj-$(CONFIG_BLK_DEV_UMEM)    += umem.o
obj-$(CONFIG_BLK_DEV_NBD)     += nbd.o
obj-$(CONFIG_BLK_DEV_CRYPTLOOP) += cryptoloop.o
obj-$(CONFIG_VIRTIO_BLK)      += virtio_blk.o

obj-$(CONFIG_BLK_DEV_SX8)     += sx8.o
obj-$(CONFIG_BLK_DEV_HD)      += hd.o

obj-$(CONFIG_XEN_BLKDEV_FRONTEND) += xen-blkfront.o
obj-$(CONFIG_XEN_BLKDEV_BACKEND)  += xen-blkback/
obj-$(CONFIG_BLK_DEV_DRBD)        += drbd/
obj-$(CONFIG_BLK_DEV_RBD)         += rbd.o
obj-$(CONFIG_BLK_DEV_PCIESSD_MTIP32XX) += mtip32xx/

obj-$(CONFIG_BLK_DEV_RSXX) += rsxx/
obj-$(CONFIG_BLK_DEV_NULL_BLK) += null_blk.o
obj-$(CONFIG_ZRAM) += zram/
obj-$(CONFIG_BLK_DEV_RD2) += rd2.o

nvme-y      := nvme-core.o nvme-scsi.o
skd-y       := skd_main.o
swim_mod-y  := swim.o swim_asm.o
```

4.4 Patch